

# Fraud Ring Radar

Detecting Coordinated Fraud Networks with Multi-Agent AI,  
Graph Memory, and Bayesian Confidence

TEAM FRAUDKILLER


NYC Tech Week · M-Agents Hackathon 2026

# Why Traditional Fraud Systems Miss Coordinated Fraud

Most fraud monitoring systems analyze transactions **one row at a time**.

## Fraud rings exploit this by:

- Staying below alert thresholds
- Splitting transfers into many small payments
- Coordinating across multiple accounts
- Sharing devices and infrastructure
- Hiding inside normal transaction activity



The fraud is not in a single transaction.

The fraud is in the graph.

# Inspiration & Vision

## Built for **Analysts** Who Only Have Minutes

We wanted to build a tool that helps fraud investigators:

- ✓ Find the suspicious ring
- ✓ Understand why it matters
- ✓ Measure confidence, not just suspicion
- ✓ Generate investigator-ready summaries

Goal: **Transform hours of investigation into minutes.**



# Live Demo Flow



## Ingest

Upload raw transaction data and device logs.



## Map

Build graph memory connecting IPs and merchants.



## Detect

Detect coordinated loops and smurfing clusters.



## Rank

Rank cases by Bayesian confidence scores.



## Co-Pilot

Ask the AI questions about case connections.



## Export

Generate investigator narratives and SAR reports.

# Multi-Agent Pipeline

## 1 GraphRAG Builder

Ingests data via **Cognee** to connect accounts, devices, and IPs.

## 2 Cluster Detector

Identifies circular money loops, smurfing, and merchant laundering hubs.

## 3 Case Summarizer

Ranks findings and produces structured SAR-style evidence reports.





# | Seeing What Rules Miss

## Traditional Rule Engine

Transaction → Threshold → Alert

Misses rings staying below limits or splitting transfers.

## Fraud Ring Radar

Graph → Relationships → Intelligence

Uncovers coordinated timing, routing hubs, and shared infrastructure.

**Benchmark Result:** The fraud ring remained below transaction thresholds but became **obvious** once viewed as a network.

# Bayesian Confidence Layer

Agents find the behavior; **PyMC** quantifies the certainty.

$$P_{\text{Fraud}} = \sum ( E_{\text{Graph}} + E_{\text{Behavior}} + E_{\text{Infra}} )$$

| Evidence Source | Signal Captured                  | Model Weighted Impact       |
|-----------------|----------------------------------|-----------------------------|
| Graph Topology  | Circular loops & Cluster size    | High (Structural certainty) |
| Behavioral Logs | Night-time bursts & reactivation | Medium (Temporal anomaly)   |
| Infrastructure  | Shared device IDs & IP clusters  | Very High (Hard linkage)    |

# | From Detection to Decision



PyMC Confidence Output:

**96%** (Interval:  
91-99%)

**⚠️ RECOMMENDED ACTION: ESCALATE**

Exposure: \$161,751 | Members: 12



# | Quantifying Business Impact

383

SUSPICIOUS FRAUD  
CASES

284

HIGH RISK ACCOUNTS

\$156K+

POTENTIAL EXPOSURE UNCOVERED

## Efficiency Gain

Complete case review in under **5 minutes**.

Traditional investigation would require hours of manual CSV filtering and visualization to trace these same loops.

# | Go-To-Market Strategy

## Target Customers



Community Banks



Credit Unions



Mid-Market FinTechs

## Privacy-First Positioning

Local deployment available. No core banking replacement required.

Explainable investigations for compliance.

## Business Model

Community Edition

Open-Source Prototype

Enterprise SaaS

Cloud-Managed Platform

On-Premise

Privacy-Sensitive Deployment

Add-Ons

Real-time Ingestion & Compliance

# | The Roadmap Ahead



## **Real-time Kafka Ingestion**

From batch processing to sub-second detection swarms.



## **Identity Resolution**

Advanced entity linking across multiple data silos.



## **SAR Automation**

One-click export for regulatory filing (FinCEN compliant).



## **Continuous Learning**

Bayesian priors that update as analysts confirm or reject alerts.



## **Core Integrations**

Native connectors for Fiserv, Jack Henry, and modern BaaS.



## **Audit Transparency**

Full history of agent actions for regulatory reviews.

# Questions?

Agents discover. PyMC quantifies. Analysts  
decide.

Thank you for your time.

Team FraudKiller | M-Agents Hackathon 2026